

MLOps: Major Assignment

Ajinkya Ghodake (G24AI1046)

GitHub Repo: [Major Assignment](#)

Step-by-Step Breakdown

- Pre-requisite for Step 1:
 - o Create local project structure

```
(base) PS D:\Projects\Study_Assignments\ML_OPS> mkdir mlops-linear-regression

Directory: D:\Projects\Study_Assignments\ML_OPS

Mode                LastWriteTime         Length Name
----                -
d-----          7/30/2025  11:15 PM             mlops-linear-regression

(base) PS D:\Projects\Study_Assignments\ML_OPS> cd .\mlops-linear-regression\

(base) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression> mkdir src, tests, .github, .github/workflows

Directory: D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression

(base) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression> New-Item -Path README.md -ItemType File
>> New-Item -Path .gitignore -ItemType File
>> New-Item -Path requirements.txt -ItemType File
>> New-Item -Path src/train.py -ItemType File -Force
>> New-Item -Path src/quantize.py -ItemType File -Force
>> New-Item -Path src/predict.py -ItemType File -Force
>> New-Item -Path src/utils.py -ItemType File -Force
>> New-Item -Path tests/test_train.py -ItemType File -Force
>> New-Item -Path .github/workflows/ci.yml -ItemType File -Force
```

Step 1: Repository Setup

- Initialize repo with:
 - o README.md

- .gitignore
- requirements.txt

```
(base) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression> git init
Initialized empty Git repository in D:/Projects/Study_Assignments/ML_OPS/mlops-linear-regression/.git/
(base) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression> echo "# mlops-linear-regression" > README.md
(base) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression> echo "__pycache__/\n*.pyc\n*.pkl\n*.joblib\n.env\n.env" > .gitignore
(base) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression> echo "scikit-learn\njoblib\n" > requirements.txt
(base) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression> git add README.md, requirements.txt, .gitignore
(base) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression> git commit -m "Step 1: Repository setup"
[master (root-commit) e85df81] Step 1: Repository setup
3 files changed, 0 insertions(+), 0 deletions(-)
create mode 100644 .gitignore
create mode 100644 README.md
create mode 100644 requirements.txt
(base) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression> git remote add origin https://github.com/git-commit-acc/mlops-linear-regression.git
(base) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression> git branch -M main
(base) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression> git push -u origin main
Enumerating objects: 5, done.
Counting objects: 100% (5/5), done.
Delta compression using up to 20 threads
Compressing objects: 100% (5/5), done.
Writing objects: 100% (5/5), 474 bytes | 474.00 KiB/s, done.
Total 5 (delta 0), reused 0 (delta 0), pack-reused 0 (from 0)
To https://github.com/git-commit-acc/mlops-linear-regression.git
 * [new branch]      main -> main
branch 'main' set up to track 'origin/main'.
(base) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression>
```

- Pre-requisite for Step 2:

- Create a conda virtual environment locally
- Install the dependencies (requirements.txt)

```
(base) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression> conda create -n mlops-linear-regression
Retrieving notices: ...working... done
Channels:
 - defaults
Platform: win-64
Collecting package metadata (repodata.json): done
Solving environment: done

## Package Plan ##

     environment location: C:\Users\ajink\anaconda3\envs\mlops-linear-regression

Proceed ([y]/n)? y

Preparing transaction: done
Verifying transaction: done
Executing transaction: done
#
# To activate this environment, use
#
#     $ conda activate mlops-linear-regression
#
# To deactivate an active environment, use
#
#     $ conda deactivate

(mlops-linear-regression) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression> pip install -r .\requirements.txt
Requirement already satisfied: scikit-learn in c:\users\ajink\appdata\local\programs\python\python313\lib\site-packages (from -r .\requirements.txt)
Requirement already satisfied: numpy in c:\users\ajink\appdata\local\programs\python\python313\lib\site-packages (from -r .\requirements.txt)
Requirement already satisfied: joblib in c:\users\ajink\appdata\local\programs\python\python313\lib\site-packages (from -r .\requirements.txt)
Collecting pytest (from -r .\requirements.txt (line 4))
  Using cached pytest-8.3.5-py3-none-any.whl (350 kB)
Installing collected packages: pytest
Successfully installed pytest-8.3.5
```

Step 2: Model Training (src/train.py)

- Load dataset.

```
src > train.py U  utils.py U X
src > utils.py > load_data
1  from sklearn.datasets import fetch_california_housing
2  from sklearn.model_selection import train_test_split
3
4  def load_data(test_size=0.2, random_state=45):
5      data = fetch_california_housing()
6      X_train, X_test, y_train, y_test = train_test_split(
7          data.data, data.target, test_size=test_size, random_state=random_state
8      )
9      return X_train, X_test, y_train, y_test
```

- Train LinearRegression model.
- Print R2 score and loss.
- Save model using joblib.

```
train.py U X  utils.py U
src > train.py > ...
1  from sklearn.datasets import fetch_california_housing
2  from sklearn.linear_model import LinearRegression
3  from sklearn.model_selection import train_test_split
4  from sklearn.metrics import mean_squared_error, r2_score
5  import joblib
6  import numpy as np
7  import sys, os
8  sys.path.append(os.path.dirname(__file__))
9  from utils import load_data
10
11 def main():
12     X_train, X_test, y_train, y_test = load_data()
13
14     model = LinearRegression()
15     model.fit(X_train, y_train)
16
17     preds = model.predict(X_test)
18     r2 = r2_score(y_test, preds)
19     mse = mean_squared_error(y_test, preds)
20
21     print(f"R2 Score: {r2:.4f}")
22     print(f"MSE: {mse:.4f}")
23
24     joblib.dump(model, "src/trained_model.joblib")
25
26 if __name__ == "__main__":
27     main()
```

- Test the code by running locally:

```
(mlops-linear-regression) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression> python .\src\train.py
R2 Score: 0.5758
MSE: 0.5559
```

- Commit changes to main branch:

```

• (mlops-linear-regression) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression> git add src/train.py, src/Utils.py, requirements.txt
• (mlops-linear-regression) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression> git commit -m "Step 2: Model Training done"
[main 5c4ff4c] Step 2: Model Training done
 3 files changed, 36 insertions(+)
 create mode 100644 src/train.py
 create mode 100644 src/Utils.py
• (mlops-linear-regression) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression> git push origin main
Enumerating objects: 8, done.
Counting objects: 100% (8/8), done.
Delta compression using up to 20 threads
Compressing objects: 100% (6/6), done.
Writing objects: 100% (6/6), 1.03 KiB | 1.03 MiB/s, done.
Total 6 (delta 0), reused 0 (delta 0), pack-reused 0 (from 0)
To https://github.com/git-commit-acc/mlops-linear-regression.git
 e85df81..5c4ff4c  main -> main
○ (mlops-linear-regression) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression>

```

Step 3: Testing Pipeline (tests/test_train.py)

- Unit test dataset loading.
- Validate model creation (LinearRegression instance).
- Check if model was trained (e.g., coef exists).
- Ensure R2 score exceeds minimum threshold.

```

test_train.py U X
tests > test_train.py > ...
1  import sys, os
2  sys.path.append(os.path.abspath(os.path.join(os.path.dirname(__file__), '..')))
3  from src.Utils import load_data
4  from sklearn.linear_model import LinearRegression
5  from sklearn.metrics import r2_score
6
7  def test_data_loading():
8      X_train, X_test, y_train, y_test = load_data()
9      assert X_train.shape[0] > 0
10     assert X_test.shape[0] > 0
11
12     def test_model_creation():
13         model = LinearRegression()
14         assert isinstance(model, LinearRegression)
15
16     def test_model_training():
17         X_train, X_test, y_train, y_test = load_data()
18         model = LinearRegression()
19         model.fit(X_train, y_train)
20         assert hasattr(model, 'coef_')
21         preds = model.predict(X_test)
22         r2 = r2_score(y_test, preds)
23         assert r2 > 0.5
24

```

- Test locally if these testcases work:

```

(mlops-linear-regression) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression> pytest tests/
===== test session starts =====
platform win32 -- Python 3.13.0, pytest-8.4.1, pluggy-1.6.0
rootdir: D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression
plugins: dash-3.0.3
collected 3 items

tests\test_train.py ... [100%]

===== 3 passed in 0.98s =====

```

Step 4: Manual Quantization (src/quantize.py)

- Load trained model.
- Extract coef and intercept .
- Save raw parameters (unquant params.joblib).
- Manually quantize them to unsigned 8-bit integers.
- Save quantized parameters (quant params.joblib).
- Perform inference with the de-quantized weights.

```

src > quantize.py > symmetric_quantize_int8
1  import joblib
2  import numpy as np
3  import sys, os
4  sys.path.append(os.path.dirname(__file__))
5  from utils import load_data
6  from sklearn.metrics import r2_score, mean_squared_error
7
8  def symmetric_quantize_int8(arr):
9      qmin, qmax = -128, 127
10
11     max_abs = np.max(np.abs(arr))
12     if max_abs == 0:
13         scale = 1.0
14         quantized = np.zeros_like(arr, dtype=np.int8)
15     else:
16         scale = max_abs / qmax
17         quantized = np.clip(np.round(arr / scale), qmin, qmax).astype(np.int8)
18
19     return quantized, scale
20
21 def symmetric_dequantize_int8(quantized, scale):
22     return quantized.astype(np.float32) * scale
23

```

```

src > quantize.py > get_size_in_kb
26
27 def symmetric_quantize_int16(arr):
28     qmin, qmax = -32768, 32767
29
30     max_abs = np.max(np.abs(arr))
31     if max_abs == 0:
32         scale = 1.0
33         quantized = np.zeros_like(arr, dtype=np.int16)
34     else:
35         scale = max_abs / qmax
36         quantized = np.clip(np.round(arr / scale), qmin, qmax).astype(np.int16)
37
38     return quantized, scale
39
40 def symmetric_dequantize_int16(quantized, scale):
41     return quantized.astype(np.float32) * scale
42
43 def get_size_in_kb(filepath):
44     size_bytes = os.path.getsize(filepath)
45     return round(size_bytes / 1024, 4)
46

```

```

src > quantize.py > symmetric_quantize_int8
40
41 def main():
42     model = joblib.load("src/trained_model.joblib")
43     coef = model.coef_
44     intercept = np.atleast_1d(model.intercept_)
45
46     joblib.dump({'coef_': coef, 'intercept_': intercept}, "src/unquant_params.joblib")
47
48     X_train, X_test, y_train, y_test = load_data()
49
50     # Evaluate original model
51     orig_preds = np.dot(X_test, coef) + intercept[0]
52     orig_r2 = r2_score(y_test, orig_preds)
53     orig_mse = mean_squared_error(y_test, orig_preds)
54
55     # Quantize model int8 parameters symmetrically
56     q_coef_int8, coef_scale_int8 = symmetric_quantize_int8(coef)
57     q_intercept_int8, int_scale_int8 = symmetric_quantize_int8(intercept)
58
59     # Quantize model int16 parameters symmetrically
60     q_coef_int16, coef_scale_int16 = symmetric_quantize_int16(coef)
61     q_intercept_int16, int_scale_int16 = symmetric_quantize_int16(intercept)
62
63     joblib.dump({
64         'q_coef': q_coef_int8,
65         'coef_scale': coef_scale_int8,
66         'q_intercept': q_intercept_int8,
67         'int_scale': int_scale_int8,
68     }, "src/quant_params_int8.joblib")
69
70     joblib.dump({
71         'q_coef': q_coef_int16,
72         'coef_scale': coef_scale_int16,
73         'q_intercept': q_intercept_int16,
74         'int_scale': int_scale_int16,
75     }, "src/quant_params_int16.joblib")
76

```

```

src > quantize.py > get_size_in_kb
47 def main():
82
83     # Paths to the saved models
84     original_model = "src/trained_model.joblib"
85     int8_model = "src/quant_params_int8.joblib"
86     int16_model = "src/quant_params_int16.joblib"
87
88     # Dequantize
89     dq_coef_int8 = symmetric_dequantize_int8(q_coef_int8, coef_scale_int8)
90     dq_intercept_int8 = symmetric_dequantize_int8(q_intercept_int8, int_scale_int8).item()
91
92     dq_coef_int16 = symmetric_dequantize_int16(q_coef_int16, coef_scale_int16)
93     dq_intercept_int16 = symmetric_dequantize_int16(q_intercept_int16, int_scale_int16).item()
94
95     # Debug
96     # print("\nSample coefficient comparison:")
97     # print(f"Original coef[:5]: {coef[:5]}")
98     # print(f"Quantized 8 bit coef[:5]: {q_coef_int8[:5]}")
99     # print(f"Dequantized 8 bit coef[:5]: {dq_coef_int8[:5]}\n")
100    # print(f"Quantized 16 bit coef[:5]: {q_coef_int16[:5]}")
101    # print(f"Dequantized 16 bit coef[:5]: {dq_coef_int16[:5]}\n")
102
103    # Evaluate quantized model
104    preds_int8 = np.dot(X_test, dq_coef_int8) + dq_intercept_int8
105    r2_int8 = r2_score(y_test, preds_int8)
106    mse_int8 = mean_squared_error(y_test, preds_int8)
107
108    preds_int16 = np.dot(X_test, dq_coef_int16) + dq_intercept_int16
109    r2_int16 = r2_score(y_test, preds_int16)
110    mse_int16 = mean_squared_error(y_test, preds_int16)
111
112    print(f"\n" + "="*50)
113    print(f"RESULTS:")
114    print(f"Original Model R2: {orig_r2:.4f}, MSE: {orig_mse:.4f}, Size: {get_size_in_kb(original_model)} KB")
115    print(f"Quantized Model 8 bit R2: {r2_int8:.4f}, MSE: {mse_int8:.4f}, Size: {get_size_in_kb(int8_model)} KB")
116    print(f"Quantized Model 16 bit R2: {r2_int16:.4f}, MSE: {mse_int16:.4f}, Size: {get_size_in_kb(int16_model)} KB")
117    print(f"\n" + "="*50)
118
119    if __name__ == "__main__":
120        main()
121

```

- Test locally to check the code:

```

=====
(mlops-linear-regression) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression> python .\src\quantize.py
=====
RESULTS:
Original Model R2: 0.5758, MSE: 0.5559, Size: 0.6807 KB
Quantized Model 8 bit R2: 0.5566, MSE: 0.5810, Size: 0.4482 KB
Quantized Model 16 bit R2: 0.5758, MSE: 0.5559, Size: 0.4492 KB
=====

```

Further Explanation on Quantization:

- 8 bit quantization was giving negative R^2 scores before (-0.1799).

#Original Logic used for Quantization:

```

def min_max_quantize(arr):
    arr_min, arr_max = arr.min(), arr.max()

```



```

    if arr_max == arr_min:
        quantized = np.full(arr.shape, 127, dtype=np.uint8)
        return quantized, arr_min, arr_max

    # Normal quantization
    quantized = ((arr - arr_min) / (arr_max - arr_min) *
255).round().astype(np.uint8)
    return quantized, arr_min, arr_max

def dequantize(quantized, arr_min, arr_max):

    if arr_max == arr_min:
        return np.full(quantized.shape, arr_min, dtype=np.float32)

    # Normal dequantization
    return quantized.astype(np.float32) / 255 * (arr_max - arr_min) + arr_min

```

Output:

```

(mlops-linear-regression) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression> python .\src\quantize.py

=====
RESULTS:
Original Model - R2: 0.575788, MSE: 0.555892
Quantized Model - R2: -0.1799, MSE: 1.5462
=====

```

- Original code failed when all values were the same. I fixed it to avoid division by zero.
- I used symmetric quantization to support both positive and negative values.
- After the code change, 8 bit quantization gave a much better R^2 than the original quantization logic (0.5566). But it is still lower than the original model R^2 value (0.5758 > 0.5566)
- Tested the quantization by switching to 16 bit for a wider range than 8 bit.
- 16 bit quantization kept values closer to original and gave better R^2 value (0.5758)

Step 5: Dockerization

Create a Dockerfile that:

- Installs dependencies
- Includes predict.py for model verification

Job Name	Description	Depends On
test_suite	Runs pytest. Must pass before others execute.	None
train_and_quantize	Trains model, runs quantization, uploads artifacts	test_suite
build_and_test_container	Builds Docker image, runs container (must execute predict.py successfully)	train_and_quantize

```

1 FROM python:3.10-slim
2 WORKDIR /app
3 COPY requirements.txt .
4 RUN pip install --no-cache-dir -r requirements.txt
5 COPY src/ src/
6 COPY tests/ tests/
7 ENTRYPOINT ["python", "src/predict.py"]

```

src/predict.py:

- Load trained model
- Run prediction on test set
- Print sample outputs

```
predict.py U X Dockerfile 1, U
src > predict.py > ...
1 import joblib
2 import sys, os
3 sys.path.append(os.path.dirname(__file__))
4 from utils import load_data
5
6 def main():
7     _, X_test, _, y_test = load_data()
8     model = joblib.load("src/trained_model.joblib")
9     preds = model.predict(X_test)
10    print("Sample predictions:", preds[:5])
11    print("Corresponding ground truths:", y_test[:5])
12
13    if __name__ == "__main__":
14        main()
15
```

```
(mlops-linear-regression) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression> python .\src\predict.py
Sample predictions: [0.71912284 1.76401657 2.70965883 2.83892593 2.60465725]
Corresponding ground truths: [0.477 0.458 5.00001 2.186 2.78 ]
```

- Test locally by building Docker image:

```
(mlops-linear-regression) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression> docker build -t mlops-lr-demo:latest .
[*] Building 30.5s (11/11) FINISHED
-> [internal] load build definition from Dockerfile
-> => transferring dockerfile: 228B
-> [internal] load metadata for docker.io/library/python:3.10-slim
-> [internal] load .dockerignore
-> => transferring context: 2B
-> [1/6] FROM docker.io/library/python:3.10-slim@sha256:81f1c3b3770d54ecfdbddcc52c2125fce674c14a1d976df8f65dc0734f9c3c5
-> => resolve docker.io/library/python:3.10-slim@sha256:81f1c3b3770d54ecfdbddcc52c2125fce674c14a1d976df8f65dc0734f9c3c5
-> [internal] load build context
-> => transferring context: 12.41kB
-> CACHED [2/6] WORKDIR /app
-> [3/6] COPY requirements.txt .
-> [4/6] RUN pip install --no-cache-dir -r requirements.txt
-> [5/6] COPY src/ src/
-> [6/6] COPY tests/ tests/
-> exporting to image
-> => exporting layers
-> => exporting manifest sha256:b183e328fcd9bdaa6b2f356fb2438c4846b1ef2e9e4e92db89107148beafcd4e
-> => exporting config sha256:988919e81aa4f07a29881ad7ec8ea7c75759b020926d6faef4853a7e5a07aa45
-> => exporting attestation manifest sha256:8ca0d430fe45e2a711209b8de2729289289c8428087f42f7ab1685ee045d37e6
-> => exporting manifest list sha256:fb3547650e3aeb1a2c7a0a7b74d8fb29c3bdabeb37d32bed2f8eal2a2150e310
-> => naming to docker.io/library/mlops-lr-demo:latest
-> => unpacking to docker.io/library/mlops-lr-demo:latest
What's next:
View a summary of image vulnerabilities and recommendations -> docker scout quickview
(mlops-linear-regression) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression>
```

```
(mlops-linear-regression) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression> docker run mlops-lr-demo:latest
Sample predictions: [0.71912284 1.76401657 2.70965883 2.83892593 2.60465725]
Corresponding ground truths: [0.477 0.458 5.00001 2.186 2.78 ]
(mlops-linear-regression) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression>
```

Step 6: CI/CD Workflow (.github/workflows/ci.yml)

Run on every push to main.

Define 3 jobs:

```
.github > workflows > ci.yml

1  name: MLOps Workflow
2
3  on:
4    push:
5      branches: [main]
6
7  jobs:
8    test-suite:
9      runs-on: ubuntu-latest
10     steps:
11       - uses: actions/checkout@v4
12       - uses: actions/setup-python@v5
13         with:
14           python-version: '3.10'
15       - run: pip install -r requirements.txt
16       - run: pytest tests/
17
18    train-and-quantize:
19      needs: test-suite
20      runs-on: ubuntu-latest
21
22     steps:
23       - uses: actions/checkout@v4
24       - uses: actions/setup-python@v5
25         with:
26           python-version: '3.10'
27
28       - name: Install dependencies
29         run: pip install -r requirements.txt
30
31       - name: Train original model
32         run: python src/train.py
33
34       - name: Run quantization (symmetric int16)
35         run: python src/quantize.py
36
37       - name: Upload artifacts
38         uses: actions/upload-artifact@v4
39         with:
40           name: model
41           path: |
42             src/trained_model.joblib
43             src/quant_params_int16.joblib
44
45    build-and-test-container:
46      needs: train-and-quantize
47      runs-on: ubuntu-latest
48      steps:
49        - uses: actions/checkout@v4
50        - run: docker build -t mlops-lr-demo .
51        - run: docker run mlops-lr-demo
```

Outputs

- Execution in local environment:

```
(mlops-linear-regression) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression> python .\src\train.py
R2 Score: 0.5758
MSE: 0.5559
(mlops-linear-regression) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression> pytest tests/
===== test session starts =====
platform win32 -- Python 3.13.0, pytest-8.4.1, pluggy-1.6.0
rootdir: D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression
plugins: dash-3.0.3
collected 3 items

tests\test_train.py ...

===== 3 passed in 0.97s =====
(mlops-linear-regression) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression> python .\src\quantize.py
=====
RESULTS:
Original Model R2: 0.5758, MSE: 0.5559, Size: 0.6807 KB
Quantized Model 8 bit R2: 0.5566, MSE: 0.5810, Size: 0.4482 KB
Quantized Model 16 bit R2: 0.5758, MSE: 0.5559, Size: 0.4492 KB
=====
(mlops-linear-regression) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression> python .\src\predict.py
Sample predictions: [0.71912284 1.76401657 2.70965883 2.83892593 2.60465725]
Corresponding ground truths: [0.477 0.458 5.00001 2.186 2.78 ]
(mlops-linear-regression) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression> |
```

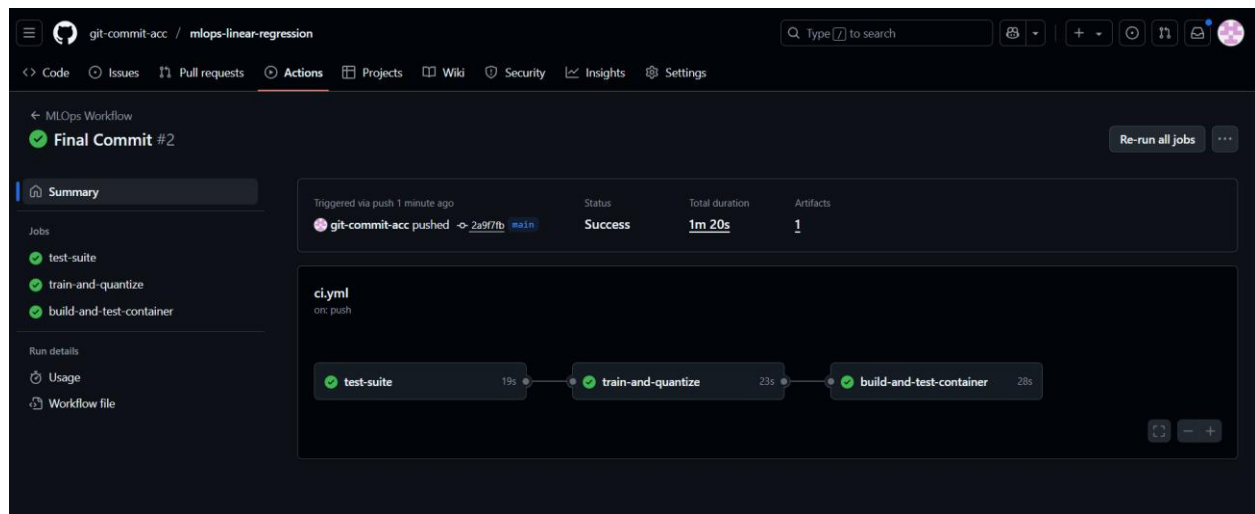
- Docker containerization:

```
(mlops-linear-regression) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression> docker build -t mlops-lr-demo:latest .

[+] Building 3.6s (11/11) FINISHED
=> [internal] load build definition from Dockerfile
=> => transferring dockerfile: 228B
=> [internal] load metadata for docker.io/library/python:3.10-slim
=> [internal] load .dockerignore
=> => transferring context: 2B
=> [1/6] FROM docker.io/library/python:3.10-slim@sha256:81f1cdb3770d54ecfdbddcc52c2125fce674c14a1d976dfd8f65dc0734f9c3c5
=> => resolve docker.io/library/python:3.10-slim@sha256:81f1cdb3770d54ecfdbddcc52c2125fce674c14a1d976dfd8f65dc0734f9c3c5
=> [internal] load build context
=> => transferring context: 2.82kB
=> CACHED [2/6] WORKDIR /app
=> CACHED [3/6] COPY requirements.txt .
=> CACHED [4/6] RUN pip install --no-cache-dir -r requirements.txt
=> CACHED [5/6] COPY src/ src/
=> CACHED [6/6] COPY tests/ tests/
=> exporting to image
=> => exporting layers
=> => exporting manifest sha256:60a3055b9f397c57025e21571c8ffe19856a8ce725f0541feda71b7ac4672b3a
=> => exporting config sha256:f484b322ce634f8c4172c5a2e1d0117d10a0d1b6e345c968699e6b343b3e0d76
=> => exporting attestation manifest sha256:acf9055f6d462f7a6326aa95d68f1c585423d285f5e60a962390a01ed535ecbe
=> => exporting manifest list sha256:6320832966fe5c1e2408a867cd4472102502e33518a053ef1115aab706d26261
=> => naming to docker.io/library/mlops-lr-demo:latest
=> => unpacking to docker.io/library/mlops-lr-demo:latest

What's next:
View a summary of image vulnerabilities and recommendations → docker scout quickview
(mlops-linear-regression) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression> docker run mlops-lr-demo:latest
Sample predictions: [0.71912284 1.76401657 2.70965883 2.83892593 2.60465725]
Corresponding ground truths: [0.477 0.458 5.00001 2.186 2.78 ]
(mlops-linear-regression) PS D:\Projects\Study_Assignments\ML_OPS\mlops-linear-regression> |
```

- Github Actions:



- Tabular comparison of all models:

Model	R ² value	MSE	Model Size
Original Model	0.5758	0.5559	0.6807 KB
8 bit Quantized (before fix)	-0.1799	1.5462	-
8 bit Quantized (after fix)	0.5566	0.5810	0.4482 KB
16 bit Quantized	0.5758	0.5559	0.4492 KB

Note: Uploading docker image to Docker hub was not mentioned in the Assignment guidelines. Hence, I have not implemented any logic to upload the Docker image to Docker hub.